
University of California, Davis

ECS145 Term Project Report

E. Kristovich, C. Nagda, V. Nguyen, V. Nguyen, R. Park

March 20, 2019

Contents

1	Introduction	1
1.1	Purpose	1
1.2	Package	1
1.3	Motivation	1
2	Implementation	1
2.1	Rcpp	1
2.2	asub()	1
2.3	acorn()	2
3	Results	3
3.1	microbenchmark	3
3.2	asub()	3
3.2.1	Case 1: Small Array	3
3.2.2	Case 2: Medium-sized Array	4
3.2.3	Case 3: High-Dimensionality Large Array	5
3.3	acorn()	6
3.3.1	Case 1: Small-sized Array	6
3.3.2	Case 2: Large Array, Simple Query	7
3.3.3	Case 3: Large Array, Complex Query	8
3.4	Case 4: acorn() using the default asub()	10
4	Analysis	10
4.1	asub()	10
4.2	acorn()	11
5	Exporting the Package	11
6	Conclusion	11
7	Contribution	11
8	Exporting Package Output	ii

1 Introduction

1.1 Purpose

The goal of this project is to mitigate slowdowns in an existing R package by implementing certain functions in C++.

1.2 Package

We have selected the `abind` package for its simplistic elegance in R, and we explore whether replacing some of the interpreted R with compiled C++ code will provide a significant speedup for arrays with many dimensions and elements. The package includes functions that allow for modifying and extracting data in n-dimensional arrays, generalizing the popular `rbind` and `cbind` functions. The `abind` package contains several functions: `abind()`, `acorn()`, `adrop()`, `afill()`, and `asub()`.

1.3 Motivation

We will implement the `asub()` and `acorn()` functions. `asub()` is commonly called as a sub-function. For example, `afill()` and `acorn()` call `asub()`. Both `asub()` and `acorn()` are frequently used functions with room for optimization because they are implemented with for loops in pure R. We believe that these two functions can benefit from speed ups by implementing them in C++, and are good targets given the scope of this project.

2 Implementation

2.1 Rcpp

We chose to use the `Rcpp` package to call our C++ functions. `Rcpp` essentially does the same thing as a `.Call()`, but it allows us to use C++ types from the `Rcpp` namespace that are much more similar to R types, making it easier to translate between R and C++.

2.2 asub()

For `asub`, we create an `asubCpp.cpp` file. As is standard, we start the `cpp` file with the following two lines to include `Rcpp`.

```
#include <Rcpp.h>
using namespace Rcpp;
```

For the function to be callable in R, we need to mark the function for export in a comment that is interpreted as an export by `Rcpp` before the function:

```
// [[Rcpp::export]]
std::string asubCpp(List idx, int numDims, NumericVector dims, bool drop, bool showDrop) {
  ...
}
```

The function itself, attached separately, replaces these lines of R code:

```

xic <- Quote(x[,drop=drop])
# Now duplicate the empty index argument the appropriate number of times
xic <- xic[c(1, 2, rep(3, length(d)), 4)]
if (is.null(drop)) {
  xic <- xic[-length(xic)]
} else {
  xic[[length(xic)]] <- drop
}
for (i in seq(along=dims))
  if (!is.null(idx[[i]]))
    xic[2+dims[i]] <- idx[i]

```

We attempted to repalce the most intensive parts of the R code: looping through the array and building a string to subset. We expect that calling a C++ function in place of the above R code will make the execution faster for arrays with upwards of 100 dimensions, as the for loop executes numDimensions times.

2.3 acorn()

For acorn(), we create an acorn1sr.cpp file. For the first few lines, it is similar with asubCpp.cpp.

```

#include <Rcpp.h>

using namespace Rcpp;

```

We also need to mark the function acorn1sr() to be exported in a comment which Rcpp will interpret as an export.

```

// [[Rcpp::export]]
List acorn1sr(int len_as, List as, NumericVector dim, int len_dim_x, List args, int len_args) {
  ...
}

```

The function acon1sr() will take length of list as, list as (this is a variable of acorn(), it it used to store the result matrix), list dim (the dimension of input matrix x), length of list dim, list args (the list of all the arguments of the acorn()), and the length of list args.

Our acorn1sr() will replace the first for loop inside function acorn.R

```

for (i in seq(4, length(dim(x)))) {
  if (length(args) >= i-3)
    a <- args[[i-3]]
  else
    a <- 1
  a <- sign(a)*min(abs(a), dim(x)[i])
  as <- c(as, list(if (a >= 0) seq(len=a) else seq(len=-a, to=dim(x)[i])))
}

```

As we expected, for loop in R code will slow down the performance. That is why we rewrote it in C++ to make it faster.

3 Results

3.1 microbenchmark

The [microbenchmark](#) package is an accurate benchmarking tool used in place of the popular *system.time*. It takes batches of sub-millisecond timings, calculates statistics such as mean and standard deviation, and is supported by *ggplot2*'s *autoplot* function to generate figures. For example, to compare two functions:

```
result <- microbenchmark(func1(), func2(), times=1000)
```

The *result* object can be printed or plotted to display the results. We use *microbenchmark* to perform experiments with 100 to 1000 trials, depending on the execution time of the experiment, and test a variety of cases. We compare our C++ implementation of functions versus the original *abind* functions.

3.2 asub()

We performed three different experiments comparing the performance of *asub()*. *abindcpp01* refers to the modified package, where as *abind* refers to the original package and acts as the control.


3.2.1 Case 1: Small Array

For our first experiment, we tested *asub()* on a trivial array with 1000 trials:

```
x <- array(1:24,dim=c(2,3,4))
asub(x, 1, 1)
```

Unit: microseconds

expr	min	lq	mean	median	uq	max	neval
abind	18.732	22.837	28.58523	28.226	30.535	125.218	1000
abindccp01	192.958	197.834	217.73057	203.222	222.466	2671.898	1000



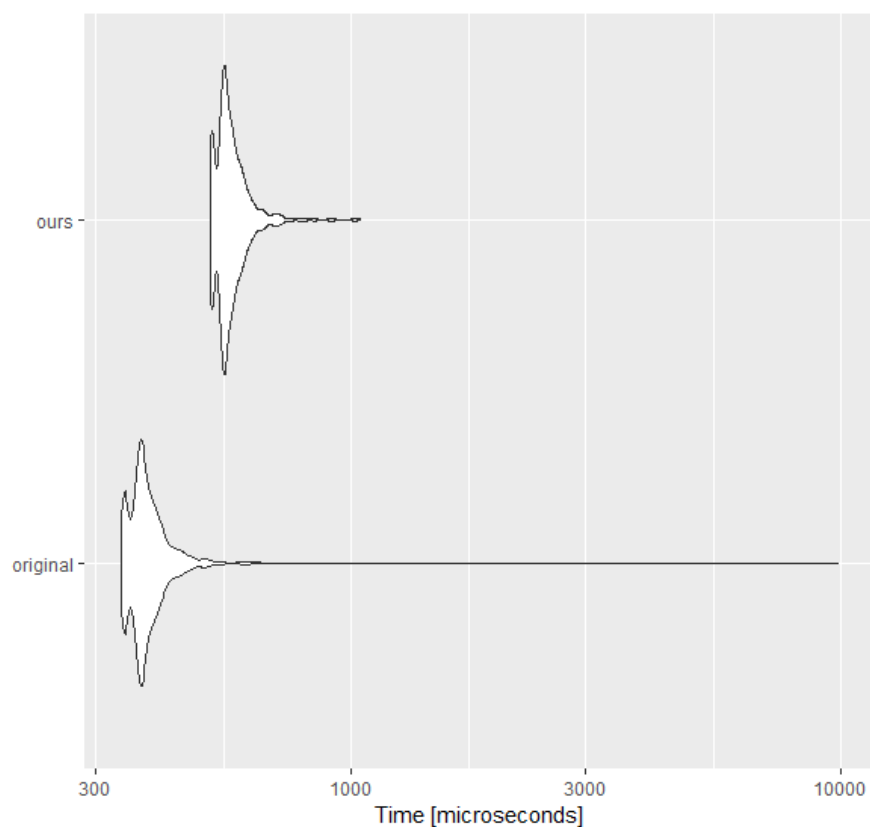
3.2.2 Case 2: Medium-sized Array

Medium-sized 6-dimension array of size 1MB, 1000 trials:

```
x <- array(1:(2**10), rep(2**3, times=6))
asub(x, 1, 1)
```

Unit: microseconds

	expr	min	lq	mean	median	uq	max	neval
	original	337.676	362.0530	393.1370	375.5235	396.3075	9797.465	1000
	ours	514.725	541.0255	569.2024	555.0100	582.8500	1038.942	1000




3.2.3 Case 3: High-Dimensionality Large Array

Large 25-dimension array of size 128MB, 100 trials:

```
x <- array(1:(2**25), rep(2, times=25))
asub(x, as.list(rep(1, 25)), 1:25)
```

Unit: milliseconds

	expr	min	lq	mean	median	uq	max	neval
original		560.0189	580.7659	611.0943	596.3183	629.2422	800.4877	100
ours		551.9873	582.6090	608.4764	592.7386	624.3100	827.2007	100



3.3 acorn()

We refer to our implementation as *abindcpp01:acorn* and the original version/control as *abind:acorn*.


3.3.1 Case 1: Small-sized Array

Small-sized array, 1000 trials:

```
x <- array(1:24,dim=c(4,3,2))
acorn(x, -3)
```

Unit: microseconds

	expr	min	lq	mean	median	uq	max	neval
	abindcpp01::acorn	317.918	326.386	360.3345	336.9065	368.8520	2697.299	1000
	abind::acorn	108.539	118.803	138.1707	127.1420	137.2775	2611.084	1000




3.3.2 Case 2: Large Array, Simple Query

Large 25-dimension array of size 128MB, a basic `asub()` call, 100 trials:

```
x <- array(1:(2**25), rep(2, times=25))
acorn(x, 1, 2, -3)
```

Unit: microseconds

expr	min	lq	mean	median	uq	max	neval
abindcpp01::acorn	489.322	498.4315	523.7388	508.8230	541.2820	647.126	100
abind::acorn	329.979	341.1400	390.8394	349.9925	373.7275	3385.223	100




3.3.3 Case 3: Large Array, Complex Query

Large 25-dimension array of size 128MB, a more complex `asub()` call, 100 trials:

```
x <- array(1:(2**25), rep(2, times=25))
acorn(x, -3, 1, 2, 4, 5, 7, 9, -3, 1, 2, 4, 5, 7, 9, -3, 1, 2, 4,
5, 7, 9, -3, 1, 2, 4, 5, 7, 9, -3, 1, 2, 4, 5, 7, 9)
```

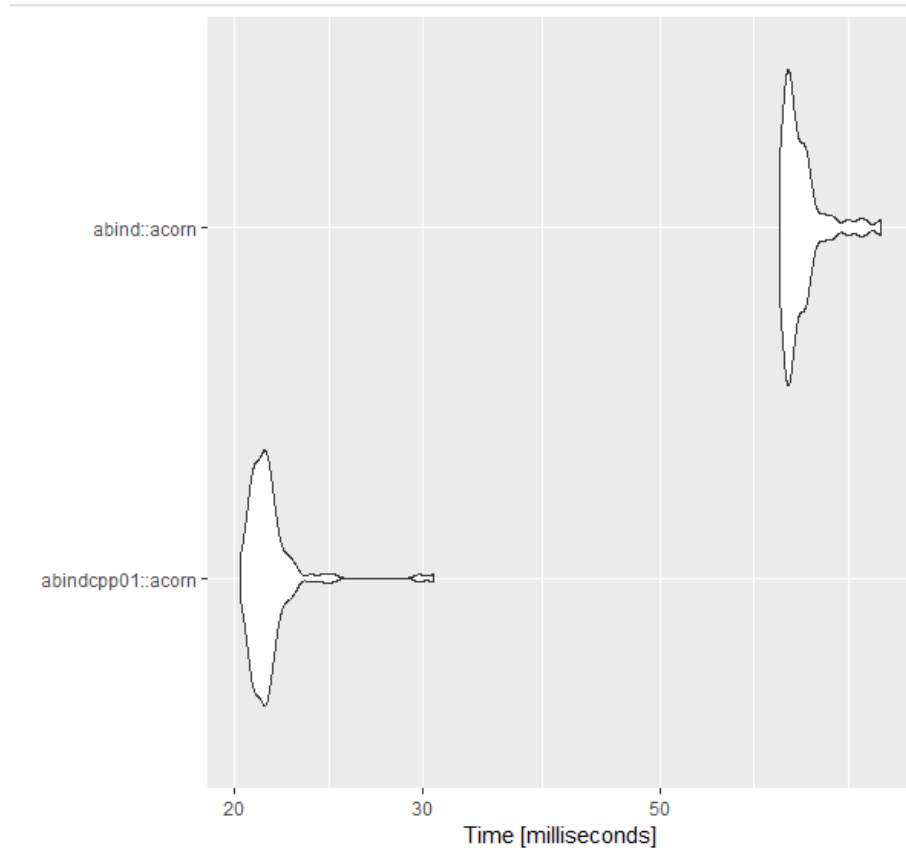
Unit: milliseconds

	expr	min	lq	mean	median	uq	max	neval
	abindcpp01::acorn	9.520858	9.811962	10.35923	10.08164	10.38878	16.52608	100
	abind::acorn	67.889450	69.604512	72.73413	70.66809	72.15222	96.52130	100



3.4 Case 4: `acorn()` using the default `asub()`

Because `acorn()` calls `asub()`, we tested our modified version of `acorn()` using the unmodified version of `asub()`. The results are of similar shape.



This is the only case where our implementation runs faster than the original.

4 Analysis

4.1 `asub()`

The `asub` function subsets arrays at indices in the specific dimensions. The original R code loops through each dimension. It constructs a string which is evaluated to subset the array. Therefore, we assumed that by looping through dimensions and constructing the string in C++, we would significantly increase performance. However, we failed to account for the time of making a C++ call. In addition to the overhead from using Rcpp the call requires passing potentially large lists and arrays to the C++ function. In the future, writing the entire function in C++ may avoid this issue. In addition, the dimensionality of an array is limited by the amount of memory R can allocate for a vector. We found that we could not construct an array with more than 25 dimensions. Thus, moving our for loop to C++ likely had little performance gain as the number of iterations would be relatively small and building a string in C++ is not much faster than building a string in R.

Despite these issues, we do see some small performance gain for high dimensionality arrays. From

our benchmarking data, we determine that when the dimension is large, we overcome time lost in the Rcpp call for an overall improvement.

4.2 acorn()

Acorn returns a small corner of an array, taking some specified slices from each dimension. It handles one, two, three, and four dimension arrays as special cases outside of any loop. After four dimensions the R code loops through each dimension and constructs a list of elements. Therefore, we assumed that by looping through dimensions and constructing the list in C++, we would significantly increase performance. This is a similar strategy to that using in asub above. However, we achieved much better results. When using a large array and a complex query with a high number of specified slices as arguments we saw close to a 10 times performance increase. We believe this is because it was more computationally intensive to construct the matrix returned by acorn than the string in asub. However, our version of the acorn function is still slower for small matrices (as seen in the graphs). This can be explained by the overhead of making the Rcpp call. Also, in the case where the dimension of a matrix x is less than 4, it will be handled by R code, making no affect to performance when using acorn's R implementation. After more testing, we could determine a threshold for the complexity of the call. Below the threshold, the original acorn would be used, and above the threshold out modified C++ acorn would be used. Thus, the performance of acorn would remain consistently high.

5 Exporting the Package

To export the package, we enter R's interactive mode and use `devtools::build()` then `devtools::check()`. The latter is analogous to R CMD check. `devtools::build()` recompiles all C++ code, and builds a tar.gz file for the package. If the `check()` commands succeeds, we can then use R CMD INSTALL `abindcpp01_1.0.tar.gz` to install the package from scratch. We copied the tests from the original abind package, so `check()` to

6 Conclusion

In conclusion, we found that we could make improvements to the abind, especially for high dimension arrays, for the asub and acorn functions. Another function, afill, calls asub and should thus also improve. Performance increased drastically in the case of acorn because the intensive list operation was shifted to a C++ function using Rcpp. Our C++ functions decreased performance in both functions for arrays with small dimensions. In the future, a threshold could be found to determine when a dimension is high enough to warrant use of an Rcpp call. This would increase the performance of the abind package overall.

7 Contribution

Reimplement asub(): E. Kristovich, C. Nagda

Reimplement acorn(): V. Nguyen, V. Nguyen

Benchmarking: R. Park

Report: all members

References

- [1] CRAN - Package abind,
<https://cran.r-project.org/web/packages/abind/index.html>

8 Exporting Package Output

```

emkristo@ad3.ucdavis.edu~$ R CMD INSTALL abindcpp01_1.0.tar.gz
* installing to library '/home/emkristo/R/x86_64-pc-linux-gnu-library/3.4'
* installing *source* package 'abindcpp01' ...
** libs
g++ -I/usr/share/R/include -DNDEBUG -I"/home/emkristo/R/x86_64-pc-linux-gnu-library/3.4/Rcpp/include" -fPIC -fPIE -pie -g -O2 -Wl,-Bsymbolic-functions -Wl,-z,relro -o abindcpp01.so RcppExports.o
installing to /home/emkristo/R/x86_64-pc-linux-gnu-library/3.4/abindcpp01/libs
** R
** inst
** preparing package for lazy loading
** help
*** installing help indices
** building package indices
** testing if installed package can be loaded
* DONE (abindcpp01)

> compileAttributes()
> compileAttributes()
> load_all()
Loading abindcpp01
Re-compiling abindcpp01
installing *source* package 'abindcpp01' ...
** libs
g++ -I/usr/share/R/include -DNDEBUG -I"/home/emkristo/R/x86_64-pc-linux-gnu-library/3.4/Rcpp/include" -fPIC -fPIE -pie -g -O2 -Wl,-Bsymbolic-functions -Wl,-z,relro -o abindcpp01.so RcppExports.o
make: Warning: File 'RcppExports.cpp' has modification time 5.5 s in the future
g++ -I/usr/share/R/include -DNDEBUG -I"/home/emkristo/R/x86_64-pc-linux-gnu-library/3.4/Rcpp/include" -fPIC -fPIE -pie -g -O2 -Wl,-Bsymbolic-functions -Wl,-z,relro -o abindcpp01.so RcppExports.o
make: warning: Clock skew detected. Your build may be incomplete.
installing to /tmp/RtmpTvm9wP/devtools_install_45bc3c2be50c/abindcpp01/libs
DONE (abindcpp01)
> build()
checking for file '/home/emkristo/ECS145/abindcpp01/DESCRIPTION' ...
preparing 'abindcpp01': (4s)
checking DESCRIPTION meta-information ...
cleaning src
checking for LF line-endings in source and make files and shell scripts

```

```

checking for empty or unneeded directories
building 'abindcpp01_1.0.tar.gz'
Warning: invalid uid value replaced by that for user 'nobody'
Warning: invalid gid value replaced by that for user 'nobody'

[1] "/home/emkristo/ECS145/abindcpp01_1.0.tar.gz"
> check()
Building abindcpp01
Setting env vars:
CFLAGS      : -Wall -pedantic -fdiagnostics-color=always
CXXFLAGS    : -Wall -pedantic -fdiagnostics-color=always
CXX11FLAGS  : -Wall -pedantic -fdiagnostics-color=always

checking for file '/home/emkristo/ECS145/abindcpp01/DESCRIPTION' ...
preparing 'abindcpp01': (2.8s)
checking DESCRIPTION meta-information ...
cleaning src
checking for LF line-endings in source and make files and shell scripts
checking for empty or unneeded directories
building 'abindcpp01_1.0.tar.gz'
Warning: invalid uid value replaced by that for user 'nobody'
Warning: invalid gid value replaced by that for user 'nobody'

Checking abindcpp01
Setting env vars:
_R_CHECK_CRAN_INCOMING_USE_ASPELL_: TRUE
_R_CHECK_CRAN_INCOMING_REMOTE_    : FALSE
_R_CHECK_CRAN_INCOMING_           : FALSE
_R_CHECK_FORCE_SUGGESTS_          : FALSE
R CMD check
using log directory '/tmp/RtmpTvm9wP/abindcpp01.Rcheck' (1.9s)
using R version 3.4.4 (2018-03-15)
using platform: x86_64-pc-linux-gnu (64-bit)
using session charset: UTF-8
using options '--no-manual --as-cran'
checking for file 'abindcpp01/DESCRIPTION'
checking extension type ... Package
this is package 'abindcpp01' version '1.0'
checking package namespace information
checking package dependencies (1.3s)
checking if this is a source package
checking if there is a namespace
checking for executable files ...
checking for hidden files and directories
checking for portable file names
checking for sufficient/correct file permissions
checking whether package 'abindcpp01' can be installed (18.2s)

```

```

checking installed package size ...
checking package directory ...
checking DESCRIPTION meta-information ...
N checking top-level files
  Non-standard file/directory found at top level:
    'core'
checking for left-over files
checking index information
checking package subdirectories ...
checking R files for non-ASCII characters ...
checking R files for syntax errors ...
checking whether the package can be loaded ...
checking whether the package can be loaded with stated dependencies ...
checking whether the package can be unloaded cleanly ...
checking whether the namespace can be loaded with stated dependencies ...
checking whether the namespace can be unloaded cleanly ...
checking loading without being on the library search path ...
checking dependencies in R code ...
checking S3 generic/method consistency (780ms)
checking replacement functions ...
checking foreign function calls ...
checking R code for possible problems (2.6s)
checking Rd files ...
checking Rd metadata ...
checking Rd line widths ...
checking Rd cross-references ...
W checking for missing documentation entries ...
  Undocumented code objects:
    'acorn1sr' 'afillcpp' 'asubCpp'
  All user-level objects in a package should have documentation entries.
  See chapter 'Writing R documentation files' in the 'Writing R
  Extensions' manual.
checking for code/documentation mismatches (711ms)
checking Rd \usage sections (884ms)
checking Rd contents ...
checking for unstated dependencies in examples ...
checking line endings in C/C++/Fortran sources/headers
checking compiled code ...
checking examples (732ms)
W checking for unstated dependencies in 'tests' ...
  'library' or 'require' call not declared from: 'abind'
checking tests ...
> > library(abind) (435ms)o 'abind.Rout.save' ...
> > library(abind) (755ms)o 'adrop.Rout.save' ...
< > (1.2s) 'afill.Rout' to 'afill.Rout.save' ...
> [1] "x" "y" (1.8s)ut' to 'asub.Rout.save' ...
> > library(abind)Rout' to 'dnns.Rout.save' ...

```

See

‘/tmp/RtmpTvm9wP/abindcpp01.Rcheck/00check.log’
for details.

R CMD check results abindcpp01 1.0
Duration: 34s

checking for missing documentation entries ... WARNING

Undocumented code objects:

‘acorn1sr’ ‘afillcpp’ ‘asubCpp’

All user-level objects in a package should have documentation entries.

See chapter ‘Writing R documentation files’ in the ‘Writing R
Extensions’ manual.

checking for unstated dependencies in ‘tests’ ... WARNING

‘library’ or ‘require’ call not declared from: ‘abind’

checking top-level files ... NOTE

Non-standard file/directory found at top level:

‘core’

0 errors | 2 warnings | 1 note